

Hamiltonian Monte Carlo and the NUTS Algorithm

MTH496 (UGP-1)
Prof. Dootika Vats

Zehaan Naik

Dept. of Mathematics and Statistics
Indian Institute of Technology Kanpur

Semester-I, 2024-25

Problem Statement

Hamiltonian Monte Carlo (HMC):

HMC uses Hamiltonian dynamics to generate distant proposals for the Metropolis algorithm.

The Main Problem:

HMC's performance depends on two parameters: step size ϵ and number of steps L . If L is too small, the algorithm behaves like a random walk; if L is too large, computation is wasted.

The NUTS Algorithm:

NUTS (as presented in Hoffman, Gelman, et al. 2014) is an extension to HMC that eliminates the need to set a number of steps L .

Metropolis Hastings Algorithm

The MH Algorithm

1. Draw $y \sim Q(x, .)$
2. Independently, draw $u \sim U(0, 1)$
3. Set $X_{n+1} = y$ if $u \leq \min(1, \alpha(x, y))$
4. Else, set $X_{n+1} = X_n$
5. Keep repeating steps 1-4 until n reaches the desired length.

Hastings's Ratio

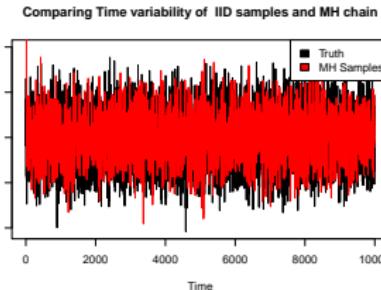
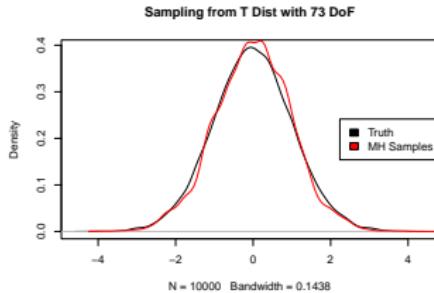
$$\alpha(x, y) = \frac{\pi(y)q(y, x)}{\pi(x)q(x, y)}$$

Metropolis Hastings Algorithm

Examples

Consider a target density $\pi \sim T_{73}$ and a proposal mechanism comprised of a Kernel $Q(x, y)$ defined as $y | x \sim N(x, 1)$. i.e.

$$q(x, y) = \frac{1}{\sqrt{2\pi}} e^{\frac{(y-x)^2}{2}}$$



The Metropolis-Hastings Algorithm

Important Results:

Theorem - 1

The MH algorithm defines the following transition kernel:

$$P(x, A) = \int_A Q(x, dy)a(x, y) + \delta_x(A) \int [1 - a(x, u)]Q(x, du).$$

Theorem - 2

The MH kernel is π -symmetric and hence π -invariant.

Hamiltonian Dynamics

Introduction:

Assume an imaginary particle a moving on a frictionless surface.
The position-momentum vector (p, q) give us:

$$U(q) = -\log(\pi(q)) \quad \text{- Potential Energy}$$

$$K(p) = \frac{p^2}{2m} \quad \text{- Kinetic Energy}$$

Together, these two quantities define the state of the particle.

Hamiltonian Dynamics

Hamiltonian Equations

The Hamiltonian is:

$$H(p, q) = K(p) + U(q),$$

And the Hamiltonian equations are:

$$\begin{aligned}\frac{dq}{dt} &= \frac{\partial H(p, q)}{\partial p} \\ \frac{dp}{dt} &= -\frac{\partial H(p, q)}{\partial q}.\end{aligned}$$

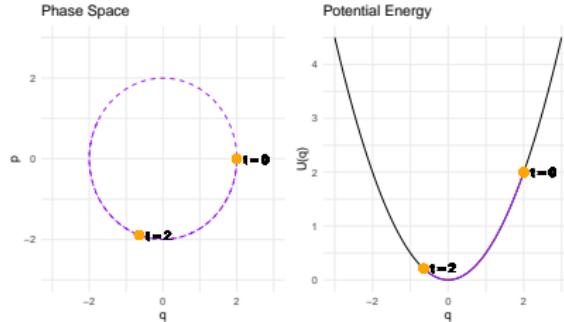
Hamiltonian Monte Carlo

Examples

Consider sampling from the target $\pi \sim N(0, 1)$ and momentum p is distributed as $N(0, 1)$. From the equations above:

$$U(q) = \frac{q^2}{2}$$

$$K(p) = \frac{p^2}{2}$$



Hamiltonian Dynamics

Assume, for a particle with state (p, q) , we get the explicit time dependence, $(p(t), q(t))$, by solving the Hamiltonian equations.

Define, $T_s(p_t, q_t) := (p_{t+s}, q_{t+s})$

Then, $T_s(T_{-s}(p_t, q_t)) = (p_t, q_t)$

An Involution

We draw motivation from the result above to define an involution,

$$P_s : (p_t, q_t, \epsilon = 1) \rightarrow (T_{\epsilon s}(p_t, q_t), \epsilon = -\epsilon).$$

Using Hamiltonian dynamics, we know,

$$|P_s| = 1.$$

Hamiltonian Monte Carlo

The HMC Algorithm

1. Sample $p \sim N(0, 1)$ and independently $u \sim U(0, 1)$
2. Set

$$(p^*, q^*) = P_s(p, q)$$

3. Set

$$\alpha(p, q) = \frac{\pi(q^*)f_P(p^*)}{\pi(q)f_P(p)}$$

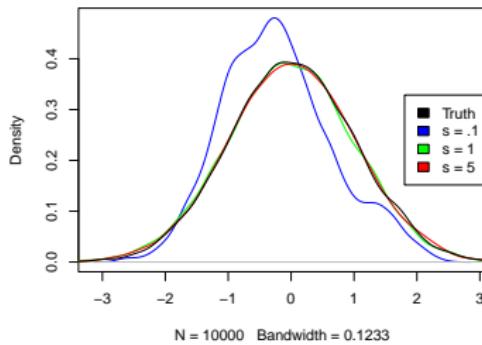
4. Set $q_{n+1} = q^*$ if $u \leq \alpha(p, q)$
5. else, Set $q_{n+1} = q$

Note: f_p is the density function for the momentum p .

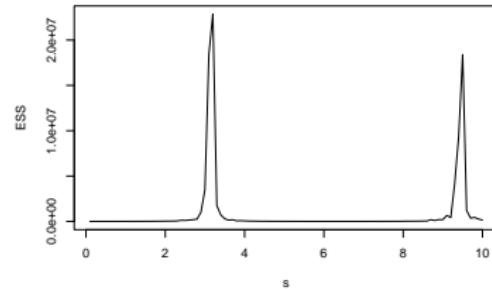
Standard Normal Example

The following plots describe the sample quality and the cyclic nature of the HMC algorithm.

Theoretical HMC



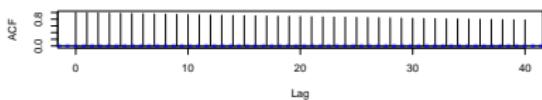
ESS values for different choices of s



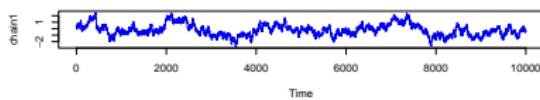
Standard Normal Example

The following plots describe the sample auto-correlation and the time variability of the HMC algorithm.

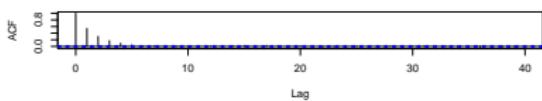
Theoretical HMC | $s = 0.1$



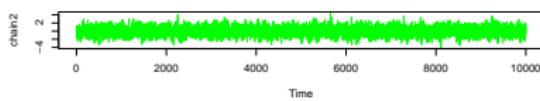
Theoretical HMC | $s = 0.1$



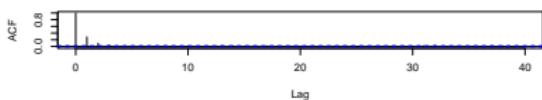
Theoretical HMC | $s = 1$



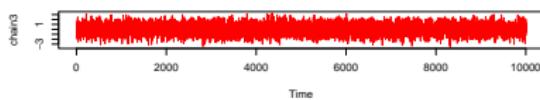
Theoretical HMC | $s = 1$



Theoretical HMC | $s = 5$



Theoretical HMC | $s = 5$



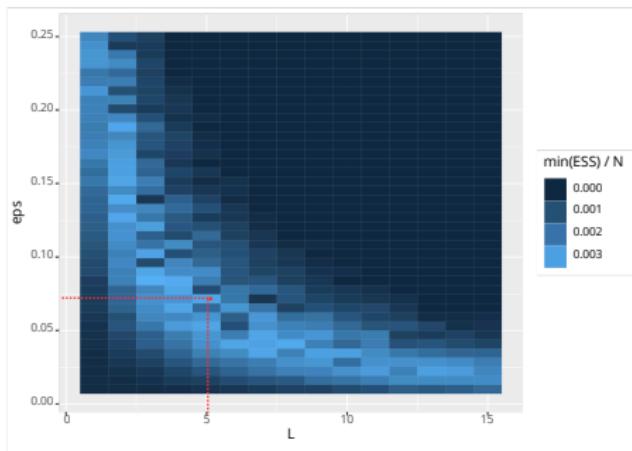
Leap Frog Integrator

Solving the Hamiltonian equations for most choices of π to obtain q_t and p_t explicitly is often not possible. Hence, we use the Leap Frog Integrator.

Note: $s = L\epsilon$

1. $p(t + \epsilon/2) = p(t) - \frac{\epsilon}{2} \frac{\partial U(q(t))}{\partial q}$
2. $q(t + \epsilon) = q(t) + \epsilon \frac{p(t + \epsilon/2)}{m}$
3. \vdots
4. $q(t + L\epsilon) = q(t + (L-1)\epsilon) + \epsilon \frac{p(t + (2L-1)\epsilon/2)}{m}$
5. $p(t + L\epsilon) = p(t + (2L-1)\epsilon/2) - \frac{\epsilon}{2} \frac{\partial U(q(t + L\epsilon))}{\partial q}$.

Optimum Choice of L and ϵ for a Multivariate Targets

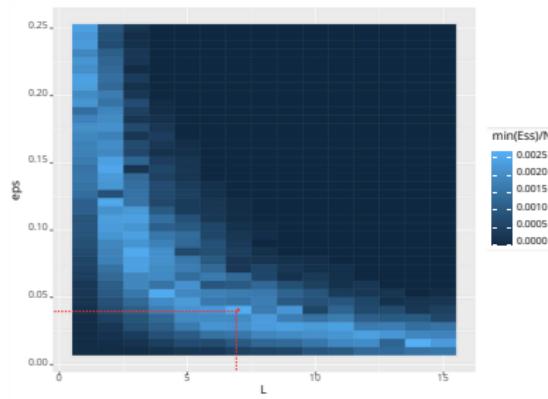


Density

Density: Multivariate Gaussian

$$\pi_G(x) = \prod_{i=1}^d \frac{1}{\sigma_i \sqrt{2\pi}} e^{-\frac{x_i^2}{2\sigma_i^2}}$$

Optimum Choice of L and ϵ for a Multivariate Targets

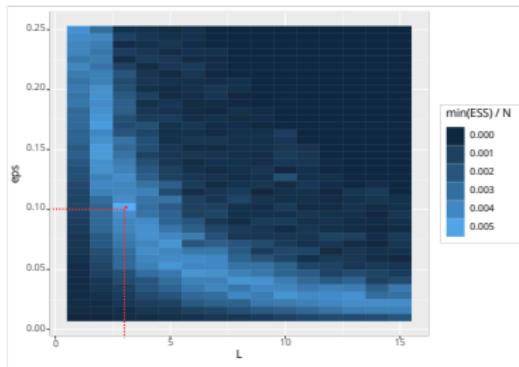


Density

Density: Logistic

$$\pi_L(x) = \prod_{i=1}^d \frac{1}{\sigma_i} \frac{e^{(x_i/\sigma_i)}}{\{1 + e^{(x_i/\sigma_i)}\}^2}$$

Optimum Choice of L and ϵ for a Multivariate Targets

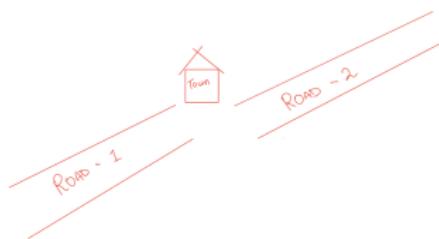


Density

Density: Skewed Gaussian

$$\pi_{SG}(x) = \prod_{i=1}^d \frac{2}{\sigma_i \sqrt{2\pi}} e^{\left(-\frac{x_i^2}{2\sigma_i^2}\right)} \Phi\left(\frac{\alpha x_i}{\sigma_i}\right)$$

NUTS: Problem Motivation



Problem Statement

- You are in a city with two roads extending in opposite directions.
- Your friend has travelled down one of these roads and got a flat tyre.
- He calls you for help but doesn't know which road he took.
- He also doesn't know how far he is from the city.
- Your task is to locate your friend and guide him back to the city.

The Naive NUTS Algorithm

Given $\theta^0, \epsilon, \mathcal{L}, M$:

for $m = 1$ to M **do**

 Resample $r^0 \sim \mathcal{N}(0, I)$.

 Resample $u \sim \text{Uniform} \left([0, \exp \left\{ \mathcal{L}(\theta^{m-1}) - \frac{1}{2} r^0 \cdot r^0 \right\}] \right)$

 Initialize $\theta^- = \theta^{m-1}$, $\theta^+ = \theta^{m-1}$, $r^- = r^0$, $r^+ = r^0$, $j = 0$, $\mathcal{C} = \{(\theta^{m-1}, r^0)\}$, $s = 1$

while $s = 1$ **do**

 Choose a direction $v_j \sim \text{Uniform}(\{-1, 1\})$

if $v_j = -1$ **then**

$\theta^-, r^-, -, -, \mathcal{C}', s' \leftarrow \text{BuildTree}(\theta^-, r^-, u, v_j, j, \epsilon)$

else

$-, -, \theta^+, r^+, \mathcal{C}', s' \leftarrow \text{BuildTree}(\theta^+, r^+, u, v_j, j, \epsilon)$

end if

if $s' = 1$ **then**

$\mathcal{C} \leftarrow \mathcal{C} \cup \mathcal{C}'$

end if

$s \leftarrow s' \mathbb{I}[(\theta^+ - \theta^-) \cdot r^- \geq 0] \mathbb{I}[(\theta^+ - \theta^-) \cdot r^+ \geq 0]$.

$j \leftarrow j + 1$

end while

 Sample θ^m, r uniformly at random from \mathcal{C}

end for

Build Tree Function

```

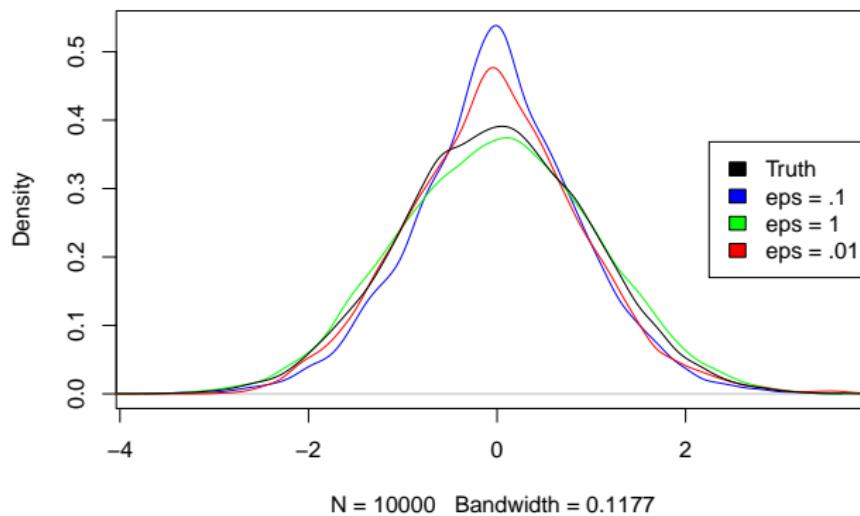
function BUILDTREE( $\theta, r, u, v, j, \epsilon$ )
  if  $j = 0$  then
    Base case: take one leapfrog step in the direction  $v$ 
     $\theta', r' \leftarrow \text{Leapfrog}(\theta, r, v\epsilon)$ 
     $C' \leftarrow \begin{cases} \{(\theta', r')\} & \text{if } u \leq \exp\left\{\mathcal{L}(\theta') - \frac{1}{2}r' \cdot r'\right\} \\ \emptyset & \text{else} \end{cases}$ 
     $s' \leftarrow \mathbb{I}\left[\mathcal{L}(\theta') - \frac{1}{2}r' \cdot r' > \log u - \Delta_{\max}\right]$ 
    return  $\theta', r', \theta', r', C', s'$ 
  else
    Recursion: build the left and right subtrees
     $\theta^-, r^-, \theta^+, r^+, C', s' \leftarrow \text{BuildTree}(\theta, r, u, v, j - 1, \epsilon)$ 
    if  $v = -1$  then
       $\theta^-, r^-, -, -, C'', s'' \leftarrow \text{BuildTree}(\theta^-, r^-, u, v, j - 1, \epsilon)$ 
    else
       $-, -, \theta^+, r^+, C'', s'' \leftarrow \text{BuildTree}(\theta^+, r^+, u, v, j - 1, \epsilon)$ 
    end if
     $s' \leftarrow s's''\mathbb{I}\left[(\theta^+ - \theta^-) \cdot r^- \geq 0\right]\mathbb{I}\left[(\theta^+ - \theta^-) \cdot r^+ \geq 0\right]$ 
     $C' \leftarrow C' \cup C''$ 
    return  $\theta^-, r^-, \theta^+, r^+, C', s'$ 
  end if
end function

```

NUTS Example

NUTS Algorithm for a $N(0, 1)$ target

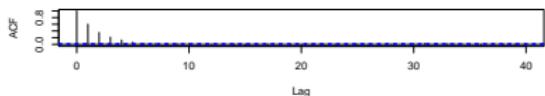
NUTS Algorithm



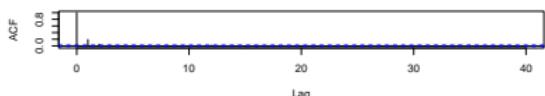
NUTS Example

Comparing different choices of ϵ

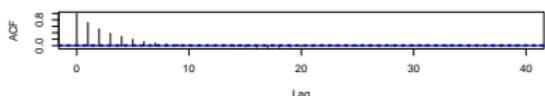
NUTS Algorithm | $\epsilon = 0.1$



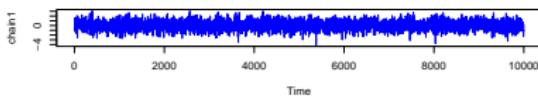
NUTS Algorithm | $\epsilon = 1$



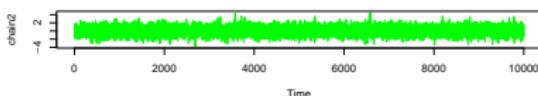
NUTS Algorithm | $\epsilon = 0.01$



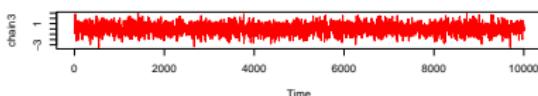
NUTS Algorithm | $\epsilon = 0.1$



NUTS Algorithm | $\epsilon = 1$



NUTS Algorithm | $\epsilon = 0.01$



Comparing NUTS and HMC

	ϵ	L_{optim}	avg(j)
1	0.0223	13	1.3813
2	0.0284	9	1.2615
3	0.0346	7	1.1897
4	0.0407	8	1.143
5	0.0469	7	1.1104
6	0.0530	5	1.0793
7	0.0592	7	1.0596
8	0.0653	4	1.0431
9	0.0715	5	1.033
10	0.0771	4	1.0244
11	0.0832	3	1.015

Table: Target: Multivariate Gaussian

	ϵ	L_{optim}	avg(j)
1	0.0223	15	1.3775
2	0.0284	10	1.2645
3	0.0346	9	1.1976
4	0.0407	8	1.1473
5	0.0469	6	1.107
6	0.0530	5	1.0791
7	0.0592	6	1.0601
8	0.0653	4	1.0423
9	0.0715	5	1.0338
10	0.0776	4	1.0239
11	0.0838	4	1.0169

Table: Target: Logistic

Future Work

I aim to explore two areas within this framework:

- **Dual Averaging:**

An extension to the NUTS Algorithm that automatically tunes ϵ . This gives us a complete problem-agnostic HMC sampler.

- **Maximising L :**

To fix the value of L based on device efficiency and maximum computation time allowed. Computing the best values for s and ϵ to both generate distant proposals and mitigate random walk behaviour.

References I

-  Hoffman, Matthew D, Andrew Gelman, et al. (2014). “The No-U-Turn sampler: adaptively setting path lengths in Hamiltonian Monte Carlo.”. In: *J. Mach. Learn. Res.* 15.1, pp. 1593–1623.
-  Holbrook, Andrew (2021). “Dr. Andrew Holbrook’s lecture on Hamiltonian Monte Carlo (HMC).”. In: *Dr. Andrew Holbrook’s lecture on Hamiltonian Monte Carlo (HMC)*.
-  Neal, Radford M (2012). *Handbook of Markov Chain Monte Carlo, Chapter 5*. arXiv preprint arXiv:1206.1901.

Thank you!